

Amazon_Food_Review-k-NN

September 25, 2018

1 K-NN on Different models-(BagOfWords,TfIDF,AvgWord2Vec,TfIDF Weighted Word2Vec)

```
In [1]: %matplotlib inline
```

```
#import all the modules
import sqlite3
import numpy as np
import pandas as pd
import nltk
import seaborn as sns

from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from sklearn.manifold import TSNE
from nltk.corpus import stopwords
from nltk.stem.porter import PorterStemmer
from nltk.stem import SnowballStemmer
from nltk.stem.wordnet import WordNetLemmatizer
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.cross_validation import cross_val_score
from collections import Counter
from sklearn import cross_validation
```

```
D:\Anaconda\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: This module v
"This module will be removed in 0.20.", DeprecationWarning)
```

```
In [2]: conn=sqlite3.connect('final2.sqlite')
        conn.cursor()
        conn.commit()
        conn.text_factory=str
        #final_data.to_sql('Reviews',conn,schema=None,if_exists='replace')
```

```
In [3]: fd=pd.read_sql_query("""SELECT * FROM REVIEWS""",conn)
```

```
In [4]: fd.head(5)
```

```
Out[4]:
```

	index	Id	ProductId	UserId	ProfileName \
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski
1	138688	150506	0006641040	A2IW4PEEK02ROU	Tracy
2	138689	150507	0006641040	A1S4A3IQ2MU7V4	sally sue "sally sue"
3	138690	150508	0006641040	AZGXZ2UUK6X	Catherine Hallberg "(Kate)"
4	138691	150509	0006641040	A3CMRKGEOP909G	Teresa

	HelpfulnessNumerator	HelpfulnessDenominator	Time \
0	0	0	939340800
1	1	1	1194739200
2	1	1	1191456000
3	1	1	1076025600
4	3	4	1018396800

	Summary \
0	EVERY book is educational
1	Love the book, miss the hard cover version
2	chicken soup with rice months
3	a good swingy rhythm for reading aloud
4	A great way to learn the months

	Text \
0	this witty little book makes my son laugh at l...
1	I grew up reading these Sendak books, and watc...
2	This is a fun way for children to learn their ...
3	This is a great little book to read aloud- it ...
4	This is a book of poetry about the months of t...

	CleanedText
0	witti littl book make son laugh loud recit car...
1	grew read sendak book watch realli rosi movi i...
2	fun way children learn month year learn poem t...
3	great littl book read nice rhythm well good re...
4	book poetri month year goe month cute littl po...

```
In [5]: conn2=sqlite3.connect('final.sqlite')
```

```
In [6]: label_df=pd.read_sql_query("""SELECT * FROM REVIEWS""",conn2)
```

```
In [7]: label_df.head(4)
```

```
Out[7]:
```

	index	Id	ProductId	UserId	ProfileName \
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski
1	138688	150506	0006641040	A2IW4PEEK02ROU	Tracy
2	138689	150507	0006641040	A1S4A3IQ2MU7V4	sally sue "sally sue"

```
3 138690 150508 0006641040 AZGXZ2UUK6X Catherine Hallberg "(Kate)"
```

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time \
0	0	0	Positive	939340800
1	1	1	Positive	1194739200
2	1	1	Positive	1191456000
3	1	1	Positive	1076025600

	Summary \
0	EVERY book is educational
1	Love the book, miss the hard cover version
2	chicken soup with rice months
3	a good swingy rhythm for reading aloud

	Text \
0	this witty little book makes my son laugh at l...
1	I grew up reading these Sendak books, and watc...
2	This is a fun way for children to learn their ...
3	This is a great little book to read aloud- it ...

	CleanedText
0	witti littl book make son laugh loud recit car...
1	grew read sendak book watch realli rosi movi i...
2	fun way children learn month year learn poem t...
3	great littl book read nice rhythm well good re...

```
In [8]: label_df=label_df.sort_values('Time',axis=0,inplace=False,kind='quicksort')
```

```
In [9]: label_df.shape
```

```
Out[9]: (364173, 12)
```

```
In [10]: fd=fd.sort_values('Time',axis=0,inplace=False,kind='quicksort')
```

```
In [11]: fd.head(3)
```

```
Out[11]:
```

	index	Id	ProductId	UserId	ProfileName \
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski
30	138683	150501	0006641040	AJ46FKX0VC7NR	Nicholas A Mesiano
424	417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina

	HelpfulnessNumerator	HelpfulnessDenominator	Time \
0	0	0	939340800
30	2	2	940809600
424	0	0	944092800

	Summary \
0	EVERY book is educational
30	This whole series is great way to spend time w...

```

424                                     Entertaining! Funny!

                                     Text \
0    this witty little book makes my son laugh at l...
30   I can remember seeing the show when it aired o...
424  Beetlejuice is a well written movie ... ever...

                                     CleanedText
0    witti littl book make son laugh loud recit car...
30   rememb see show air televis year ago child sis...
424  beetlejuic well written movi everyth excel act...

```

2 Since we have not much of RAM we are working on select set of samples

```
In [12]: #work on selected samples
```

```

n_samples=2000
test_data=fd.sample(n_samples)    ##using pd.sample() function
label_data=label_df['Score'][0:2000]

```

```
In [13]: test_data.shape
```

```
Out[13]: (2000, 11)
```

```
In [14]: label_data.shape
```

```
Out[14]: (2000,)
```

3 Bag Of Words KNN

```
In [15]: #Bag of Words
```

```

count_vect=CountVectorizer()
final_count=count_vect.fit_transform(test_data['CleanedText']).values
print("the type of count vectorizer is:",type(final_count))
final_count.get_shape()

```

```
the type of count vectorizer is: <class 'scipy.sparse.csr.csr_matrix'>
```

```
Out[15]: (2000, 6910)
```

```
In [16]: #split the data into train and test fo bag of words
```

```

X_train,X_test,Y_train,Y_test=cross_validation.train_test_split(final_count,label_data)
#split train into cross val train and cross val test
X_t,X_cv,Y_t,Y_cv=cross_validation.train_test_split(X_train,Y_train,test_size=0.3)

```

```

In [17]: #shape of train and test,cv
         X_t.shape

Out[17]: (980, 6910)

In [18]: X_cv.shape

Out[18]: (420, 6910)

In [19]: X_test.shape

Out[19]: (600, 6910)

In [20]: Y_test.shape

Out[20]: (600,)

In [21]: X_train.shape

Out[21]: (1400, 6910)

In [22]: Y_train.shape

Out[22]: (1400,)

In [23]: #find the best k based on cv accuracy for bow

         for i in range(1,10,2):
             # instantiate learning model (k = 10)
             knn = KNeighborsClassifier(n_neighbors=i)

             # fitting the model on crossvalidation train
             knn.fit(X_t, Y_t)

             # predict the response on the crossvalidation train
             pred = knn.predict(X_cv)

             # evaluate CV accuracy
             acc = accuracy_score(Y_cv, pred, normalize=True) * float(100)
             print('\nCV accuracy for k = %d is %d%%' % (i, acc))

         #test accuracy
         knn = KNeighborsClassifier(7)
         knn.fit(X_t,Y_t)
         pred = knn.predict(X_test)
         acc = accuracy_score(Y_test, pred, normalize=True) * float(100)
         print('\n****Test accuracy for k = 7 is %d%%' % (acc))

```

CV accuracy for k = 1 is 77%

CV accuracy for k = 3 is 85%

CV accuracy for k = 5 is 87%

CV accuracy for k = 7 is 88%

CV accuracy for k = 9 is 88%

****Test accuracy for k = 7 is 91%

In [24]: *#with either k=8 or 9 we get best accuracy for test to be 91%,now with 10 fold cross*

```
n_list=list(range(0,30))
neighb=list(filter(lambda x: x % 2 != 0, n_list))
#create a list of cross-val scores
scores_cv=[]
for k in neighb:
    knn=KNeighborsClassifier(n_neighbors=k)
    scores=cross_val_score(knn,X_train,Y_train,cv=10,scoring='accuracy')
    scores_cv.append(scores.mean())
```

In [25]: *#Find misclassification error(i.e)how much data is misclassified*

```
MSE=[1-x for x in scores_cv]

#find optimal k

optimal_k=neighb[MSE.index(min(MSE))]
print('\n the optimal k is %d.' % optimal_k)
```

the optimal k is 9.

In [26]: *#KNN with optimal k and test accuracy for bag of words*

```
knn_opt=KNeighborsClassifier(n_neighbors=optimal_k)
#fit the model
knn_opt.fit(X_train,Y_train)
#predict the model
prediction=knn_opt.predict(X_test)

#the accuracy score
acc_score=accuracy_score(Y_test,pred)* 100
print('\n the accuracy score for bag of words model with optimal k=%d is %f%%' %(optimal_k,acc_score))
```

the accuracy score for bag of words model with optimal k=9 is 91.500000%

4 TF-IDF KNN

```
In [27]: tf_idf_vect=TfidfVectorizer()
         final_tf_idf_vect=tf_idf_vect.fit_transform(test_data["CleanedText"].values)
         final_tf_idf_vect.get_shape()
```

```
Out[27]: (2000, 6910)
```

```
In [28]: #split the data into train and test fo tf-idf
```

```
         X_train,X_test,Y_train,Y_test=cross_validation.train_test_split(final_tf_idf_vect,labels,
         #split train into cross val train and cross val test
         X_t,X_cv,Y_t,Y_cv=cross_validation.train_test_split(X_train,Y_train,test_size=0.3)
```

```
In [29]: X_train.shape
```

```
Out[29]: (1400, 6910)
```

```
In [30]: X_test.shape
```

```
Out[30]: (600, 6910)
```

```
In [31]: X_t.shape
```

```
Out[31]: (980, 6910)
```

```
In [32]: Y_t.shape
```

```
Out[32]: (980,)
```

```
In [33]: X_cv.shape
```

```
Out[33]: (420, 6910)
```

```
In [34]: #find the best k based on cv accuracy for bow
```

```
         for i in range(1,10,2):
             # instantiate learning model (k = 10)
             knn = KNeighborsClassifier(n_neighbors=i)

             # fitting the model on crossvalidation train
             knn.fit(X_t, Y_t)

             # predict the response on the crossvalidation train
             pred = knn.predict(X_cv)

             # evaluate CV accuracy
             acc = accuracy_score(Y_cv, pred, normalize=True) * float(100)
             print('\nCV accuracy for k = %d is %d%%' % (i, acc))
         #test accuracy
```

```

knn = KNeighborsClassifier(9)
knn.fit(X_t,Y_t)
pred = knn.predict(X_test)
acc = accuracy_score(Y_test, pred, normalize=True) * float(100)
print('\n****Test accuracy for k = 9 is %d%%' % (acc))

```

CV accuracy for k = 1 is 79%

CV accuracy for k = 3 is 83%

CV accuracy for k = 5 is 84%

CV accuracy for k = 7 is 86%

CV accuracy for k = 9 is 86%

****Test accuracy for k = 9 is 91%

```

In [35]: #for finding optimal k with odd list for neighbors
n_list=list(range(0,30))
neighb=list(filter(lambda x: x % 2 != 0, n_list))
#create a list of cross-val scores
scores_cv=[]
for k in neighb:
    knn=KNeighborsClassifier(n_neighbors=k)
    scores=cross_val_score(knn,X_train,Y_train,cv=10,scoring='accuracy')
    scores_cv.append(scores.mean())

```

```

In [36]: #Find misclassification error(i.e)how much data is misclassified

```

```

MSE=[1-x for x in scores_cv]

#find optimal k

optimal_k=neighb[MSE.index(min(MSE))]
print('\n the optimal k is %d.' % optimal_k)

```

the optimal k is 15.

```

In [37]: #KNN with optimal k and test accuracy for tf_idf model

```

```

knn_opt=KNeighborsClassifier(n_neighbors=optimal_k)
#fit the model
knn_opt.fit(X_train,Y_train)
#predict the model

```



```

prediction=knn_opt.predict(X_test)

#the accuracy score
acc_score=accuracy_score(Y_test,pred)* 100
print('\n the accuracy score for bag of words model with optimal k=%d is %f%%' %(optimal_k, acc_score))

```

the accuracy score for bag of words model with optimal k=15 is 91.166667%

5 AvgWord2Vec KNN

```

In [38]: import gensim
         from gensim.models import word2vec, KeyedVectors

```

D:\Anaconda\lib\site-packages\gensim\utils.py:1209: UserWarning: detected Windows; aliasing chunkize to chunkize_serial
 warnings.warn("detected Windows; aliasing chunkize to chunkize_serial")

```

In [39]: model=word2vec.load('w2vmodel.model')

```

```

-----

AttributeError                                Traceback (most recent call last)

<ipython-input-39-ff5ebc47b102> in <module>()
----> 1 model=word2vec.load('w2vmodel.model')

```

AttributeError: module 'gensim.models.word2vec' has no attribute 'load'

```

In [40]: #word2vec own model
         i=0
         list_of_sentence=[]
         for sent in test_data['CleanedText'].values:
             list_of_sentence.append(sent.split())
             #sent=cleanhtml(sent)
             #for w in sent.split():
                 #for cleaned in cleanpunc(w).split():
                     #if(cleaned.isalpha()):
                         #filtered_sentence.append(cleaned.lower())
                     #else:
                         #continue
             #list_of_sentence.append(filtered_sentence)
         print(test_data['CleanedText'].values[0])
         print('#####')

```

```

print(list_of_sentence[0])
w2v_model=gensim.models.Word2Vec(list_of_sentence,min_count=5,size=50,workers=4)

words=list(w2v_model.wv.vocab)
print(len(words))

```

sorri food snob allergi dont spit mouth finish chew mean tast good certain doesnt make graham

['sorri', 'food', 'snob', 'allergi', 'dont', 'spit', 'mouth', 'finish', 'chew', 'mean', 'tast'
1990

```

In [41]: sent_vectors = []
        for sent in list_of_sentence: # for each review/sentence
            sent_vec = np.zeros(50) # as word vectors are of zero length
            cnt_words =0 # num of words with a valid vector in the sentence/review
            for word in sent: # for each word in a review/sentence
                if word in words:
                    vec = w2v_model.wv[word]
                    sent_vec += vec
                    cnt_words += 1
            if cnt_words != 0:
                sent_vec /= cnt_words
            sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))

vec_avg=np.array(sent_vectors)

```

2000
50

```

In [42]: #split the data into train and test fo tf-idf

X_train,X_test,Y_train,Y_test=cross_validation.train_test_split(vec_avg,label_data,tes
#split train into cross val train and cross val test
X_t,X_cv,Y_t,Y_cv=cross_validation.train_test_split(X_train,Y_train,test_size=0.3)

```

```

In [44]: X_train.shape

```

```

Out[44]: (1400, 50)

```

```

In [45]: Y_cv.shape

```

```

Out[45]: (420,)

```

```

In [46]: Y_train.shape

```

```

Out[46]: (1400,)

```

```

In [47]: X_test.shape
Out[47]: (600, 50)

In [48]: X_cv.shape
Out[48]: (420, 50)

In [49]: X_t.shape
Out[49]: (980, 50)

In [53]: #find the best k based on cv accuracy for AvgWord2Vec

for i in range(1,10,2):
    # instantiate learning model (k = 10)
    knn = KNeighborsClassifier(n_neighbors=i)

    # fitting the model on crossvalidation train
    knn.fit(X_t, Y_t)

    # predict the response on the crossvalidation train
    pred = knn.predict(X_cv)

    # evaluate CV accuracy
    acc = accuracy_score(Y_cv, pred, normalize=True) * float(100)
    print('\nCV accuracy for k = %d is %d%%' % (i, acc))
#test accuracy
knn = KNeighborsClassifier(9)
knn.fit(X_t,Y_t)
pred = knn.predict(X_test)
acc = accuracy_score(Y_test, pred, normalize=True) * float(100)
print('\n****Test accuracy for k = 9 is %d%%' % (acc))

CV accuracy for k = 1 is 80%

CV accuracy for k = 3 is 85%

CV accuracy for k = 5 is 87%

CV accuracy for k = 7 is 88%

CV accuracy for k = 9 is 89%

****Test accuracy for k = 9 is 91%

In [54]: #for finding optimal k with odd list for neighbors
n_list=list(range(0,30))

```

```

neighb=list(filter(lambda x: x % 2 != 0, n_list))
#create a list of cross-val scores
scores_cv=[]
for k in neighb:
    knn=KNeighborsClassifier(n_neighbors=k)
    scores=cross_val_score(knn,X_train,Y_train,cv=10,scoring='accuracy')
    scores_cv.append(scores.mean())

```

In [55]: *#Find misclassification error(i.e)how much data is misclassified*

```

MSE=[1-x for x in scores_cv]

#find optimal k

optimal_k=neighb[MSE.index(min(MSE))]
print('\n the optimal k is %d.' % optimal_k)

```

the optimal k is 9.

In [57]: *#KNN with optimal k and test accuracy for AvgWord2Vec model*

```

knn_opt=KNeighborsClassifier(n_neighbors=optimal_k)
#fit the model
knn_opt.fit(X_train,Y_train)
#predict the model
prediction=knn_opt.predict(X_test)

#the accuracy score
acc_score=accuracy_score(Y_test,pred)* 100
print('\n the accuracy score for AvgWord2Vec model with optimal k=%d is %f%%' %(optimal_k,acc_score))

```

the accuracy score for AvgWord2Vec model with optimal k=9 is 91.166667%

6 Weighted Tf-Idf Word2Vec KNN

In [58]: *# TF-IDF weighted Word2Vec*

```

tf_idf_features = tf_idf_vect.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = [] # the tfidf-w2v for each sentence/review is stored in this list
row=0
for sent in list_of_sentence: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0 # num of words with a valid vector in the sentence/review

```

```

for word in sent: # for each word in a review/sentence
    if word in words:
        vec = w2v_model.wv[word]
        # obtain the tf-idfidf of a word in a sentence/review
        tf_idf = final_tf_idf_vect[row, tf_idf_features.index(word)]
        sent_vec += (vec * tf_idf)
        weight_sum += tf_idf
if weight_sum != 0:
    sent_vec /= weight_sum
tfidf_sent_vectors.append(sent_vec)
row += 1

```

In [59]: `tf_vec_avg=np.array(tfidf_sent_vectors)`

In [60]: *#split the data into train and test fo tf-idf*

```

X_train,X_test,Y_train,Y_test=cross_validation.train_test_split(tf_vec_avg,label_data
#split train into cross val train and cross val test
X_t,X_cv,Y_t,Y_cv=cross_validation.train_test_split(X_train,Y_train,test_size=0.3)

```

In [62]: `X_train.shape`

Out[62]: (1400, 50)

In [63]: `X_t.shape`

Out[63]: (980, 50)

In [64]: `X_cv.shape`

Out[64]: (420, 50)

In [65]: `X_test.shape`

Out[65]: (600, 50)

In [66]: *#find the best k based on cv accuracy for TfIDF Weighted Word2Vec*

```

for i in range(1,10,2):
    # instantiate learning model (k = 10)
    knn = KNeighborsClassifier(n_neighbors=i

    # fitting the model on crossvalidation train
    knn.fit(X_t, Y_t)

    # predict the response on the crossvalidation train
    pred = knn.predict(X_cv)

    # evaluate CV accuracy
    acc = accuracy_score(Y_cv, pred, normalize=True) * float(100)

```

```

        print('\nCV accuracy for k = %d is %d%%' % (i, acc))
    #test accuracy
    knn = KNeighborsClassifier(9)
    knn.fit(X_t,Y_t)
    pred = knn.predict(X_test)
    acc = accuracy_score(Y_test, pred, normalize=True) * float(100)
    print('\n****Test accuracy for k = 9 is %d%%' % (acc))

```

CV accuracy for k = 1 is 81%

CV accuracy for k = 3 is 84%

CV accuracy for k = 5 is 85%

CV accuracy for k = 7 is 85%

CV accuracy for k = 9 is 85%

****Test accuracy for k = 9 is 91%

```

In [67]: #for finding optimal k with odd list for neighbors
n_list=list(range(0,30))
neighb=list(filter(lambda x: x % 2 != 0, n_list))
#create a list of cross-val scores
scores_cv=[]
for k in neighb:
    knn=KNeighborsClassifier(n_neighbors=k)
    scores=cross_val_score(knn,X_train,Y_train,cv=10,scoring='accuracy')
    scores_cv.append(scores.mean())

```

```

In [68]: #Find misclassification error(i.e)how much data is misclassified

```

```

MSE=[1-x for x in scores_cv]

#find optimal k

optimal_k=neighb[MSE.index(min(MSE))]
print('\n the optimal k is %d.' % optimal_k)

```

the optimal k is 15.

```

In [69]: #KNN with optimal k and test accuracy for Weighted TfIDF Word2Vec model

```

```

knn_opt=KNeighborsClassifier(n_neighbors=optimal_k)
#fit the model

```

```

knn_opt.fit(X_train,Y_train)
#predict the model
prediction=knn_opt.predict(X_test)

#the accuracy score
acc_score=accuracy_score(Y_test,pred)* 100
print('\n the accuracy score for TfIdf Word2Vec model with optimal k=%d is %f%%' %(op

```

the accuracy score for TfIdf Word2Vec model with optimal k=15 is 91.166667%

Conclusion/Observation

With optimal k being 9 or 15 we get a accuracy of 91% and 91.1666%